

EECS 341 Spring 209

Introduction to Databases

Instructor: Éamon Johnson, ebj8 @ case, Olin 504

Classroom: Rockefeller 301 T-TH 2:30 - 3:45 pm

Web page: on Canvas (<https://canvas.case.edu/>)

Office hours: TBD

TAs: TBD

Course Style

- **Basic Material Course**

- **Project-oriented course**

You can learn a lot of project-related stuff in this course—or, a lot less depending on how much time you put into it.

There have been excellent projects completed in this course in the past!

- Start installing your DBMS (MySQL) right away.

- Start installing Virtual Box with Ubuntu right away.

- Start learning about web-based application development.

Textbook, DBMS, Exams

Text: “Database Systems Concepts”,
A. Silberschatz, H. Korth, and S. Sudarshan,
6th edition 2011.


Project DBMS: MySQL

Midterm Exam: Tuesday, March 5, 2019, in class

Final Exam: Thursday, April 25, 2019, in class

Quizzes: Dates to be announced.

Project Presentations: May 7, 2019



Grading

- Assignments (~6) 20%
 - Project 30%
 - Quizzes (2) 10%
 - Midterm exam 20%
 - Final exam 20%
-
- **Late Assignment and Project Policy:** All homeworks are due at noon on the due date. Late submissions are accepted until the midnight of the due date subject to 50% penalty.
-
- **Makeup Exams:** A priori notice for any exams missed is needed. Makeup exams will NOT be given except with a serious, documented medical or legal excuse. There will not be any makeup quizzes.

Course Objectives

- In-depth understanding of basic concepts in database systems including database modeling, query languages, storage structures and query optimization issues.
- In-depth knowledge of the Entity Relationship Model and the Relational Model, as well as steps in designing, populating and maintaining a database.
- In-depth knowledge of querying relational databases, and query languages, including relational algebra, relational calculus, and SQL (with aggregate functions).
- General knowledge of query processing and transaction management in databases.
- In-depth knowledge of logical database design, schema enhancement and normalization using data dependencies.

Course Objectives (2)

- Hands-on experience on building a database-enabled web-based application that involves:
 - Designing and querying a relational database,
 - Using a (commercial MySQL) database management system effectively,
 - Project: Building a web application with a database backend.
 - Building a basic user interface for the course project.

- Experience in
 - Report writing (project proposal, progress report and final project report are required),
 - Project design, demonstration and presentation,
 - Experience in working in a project team

Course Prerequisites

- Substantial experience with at least one programming language.
EECS 233: Intro to Data Structures is a prerequisite to this course.
- Knowledge of basic data structures including trees, linked lists, arrays, priority queues, heaps.
The course covers (somewhat) B+-trees.
- Basic knowledge of Hashing, Searching and Sorting.
The course will use and refer to these topics without any overview--in the query processing part of the course.
- Knowledge of first-order logic (covered in EECS 302: Discrete math).
This knowledge is needed for relational calculus coverage.
- Mathematical maturity with set theory, functions, notations, symbolisms, and algorithmic thinking.

Course Topics—First Half

- Overview of Database Systems
- Database Design-- Entity-Relationship Data Model
- Relational Data Model
- Logical Database Design, ER to Relational and Views
- Relational Algebra
- Relational Calculus
- SQL: The Database Language

Course Topics—Second Half

- Web-Enabled Database Application Development
- Relational Database Query Processing Techniques
- Dependency Theory: Relational Database Design, Data Dependencies, Normal Forms, Schema Refinement
- Transaction Management Basics
- Data Warehouses/ OLAP

What Do We *Really* Learn in This Course?

- **Database models, query languages, query processing, transaction management**

Very important practical material!


Competency is tested in assignments and exams.

- **Project Building; Team Work!**
- **Report Writing**
- **A lot of informally (yet reasonably precisely) expressed concepts/terminology.**

How to do well

- Be resourceful!
- Install a database immediately (MySQL)
- Read the book before class, then come to class to discuss the material
- Go to office hours
- Talk to your peers
- Be honest: all words and code must be your own
 - No notes from peer discussions
 - No copying from the internet

Why do we need a database?



- Find me a book
 - Give me some time ...
- We have what you need
 - Wait, it was gone....
- I can't find the book
 - Maybe someone took it without telling me...




Databases: Convenient and Efficient



- Find me a book
 - Here you go!
- We have what you need
 - Happy to serve!
- I can't find the book
 - Wait for the new order



What is a Database?



- A very large, integrated collection of bits.
- Models real-world enterprise.
 - Entities (e.g., students, courses)
 - Relationships (e.g., student X is taking EECS341)
- A Database Management System (DBMS) is a software package designed to store and manage databases.

Files vs. Databases

- Application must stage large datasets between main memory and secondary storage (e.g., buffering, page-oriented access, 32-bit addressing, etc.)
- Special code for different queries
- Must protect data from inconsistency due to multiple concurrent users
- Crash recovery
- Security and access control

Major Functionality of Database

- Data Independence
 - No need for knowing physical schema
- Efficient access
 - Query language and execution optimization
- Reduced application development time
 - Unified data format, APIs, ...
- Data integrity and security
 - Role-based access control
- Uniform data administration
- Concurrent access, recovery from crashes



Why Study Databases??

- Shift from computation to information
 - at the “low end”: dynamic web spaces
 - at the “high end”: scientific applications
- Datasets increasing in diversity and volume.
 - Digital libraries, interactive video, Human Genome project, Earth-Observing Satellite (EOS) project
 - the need for DBMS is exploding
- DBMS encompasses most of CS
 - OS, languages, theory, AI, multimedia, logic

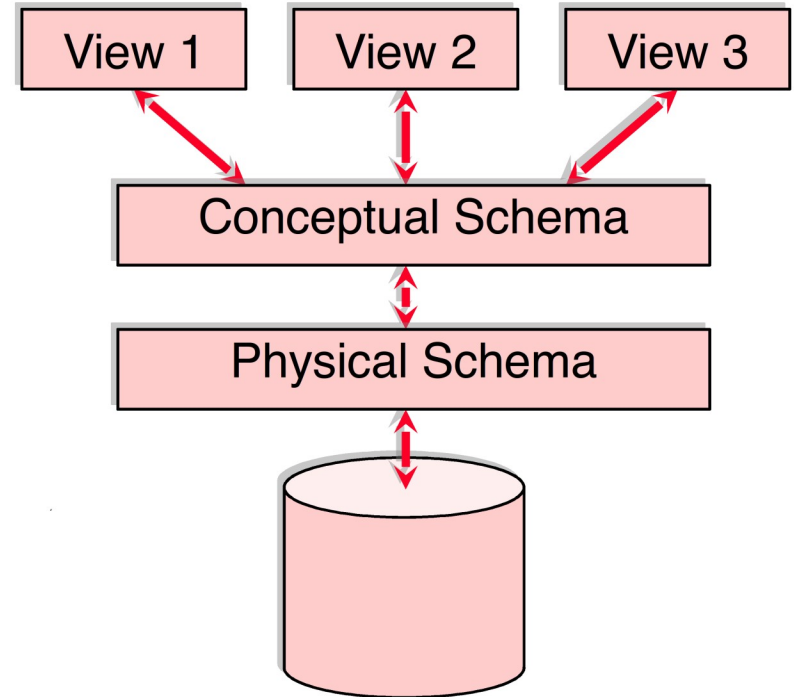


Data Models

- A *data model* is a collection of concepts for describing data.
- A *schema* is a description of a particular collection of data, using the a given data model.
- The *relational model of data* is the most widely used model today.
 - Main concept: *relation*, basically a table with rows and columns.
 - Every relation has a *schema*, which describes the allowed contents of columns, or fields.

Levels of Abstraction

- Many views, single conceptual (logical) schema and physical schema.
 - Views describe how users see the data.
 - Conceptual schema defines logical structure
 - Physical schema describes the files and indexes used.



*Schemas are defined using a Data-Description Languages (DDLs)
Data is modified/queried using Data-Management Languages (DMLs).*

Example: University Database

- Conceptual schema:
 - *Students(sid: string, name: string, login: string, dob: date, gpa: real)*
 - *Courses(cid: string, cname: string, credits: integer)*
 - *Enrolled(sid: string, cid: string, grade: string)*
- Physical schema:
 - Relations stored as unordered files.
 - Index on first column of Students.
- External Schema (View):
 - *Course_info(cid: string, enrollment: integer)*

Data Independence*

- Applications insulated from how data is actually structured and stored.
- Logical data independence: Protection from changes in *logical* structure of data.
- Physical data independence: Protection from changes in *physical* structure of data.

☛ *One of the most important benefits of using a DBMS!*

Concurrency Control

- Concurrent execution of multiple user queries is essential for good DBMS performance.
 - Because disk accesses are frequent, and relatively slow, it is important to keep the cpu humming by working on several user programs concurrently.
- Interleaving actions of different user programs can lead to inconsistency: e.g., check is cleared while account balance is being computed.
- DBMS ensures such problems don't arise: users can pretend they are using a single-user system.

Database Transactions

- Key concept is transaction (Xact), which is an *atomic* sequence of database actions.
- Each transaction, executed completely, must leave the DB in a consistent state if DB is consistent when the transaction begins.
 - Users can specify some simple integrity constraints on the data, and the DBMS will enforce these constraints.
 - Beyond this, the DBMS does not really understand the semantics of the data. (e.g., it does not understand how the interest on a bank account is computed).
 - Thus, ensuring that a transaction (run alone) preserves consistency is ultimately the *user's* responsibility!


Scheduling Concurrent Transactions

- DBMS ensures that execution of $\{T_1, \dots, T_n\}$ is equivalent to some serial execution $T_1' \dots T_n'$.
 - Before reading/writing an object, a transaction requests a lock on the object, and waits till the DBMS gives it the lock. All locks are released at the end of the transaction. (Strict Two-Phase Locking (2PL) protocol.)
 - **Idea:** If an action of T_i (say, writing X) affects T_j (which perhaps reads X), one of them, say T_i , will obtain the lock on X first and T_j is forced to wait until T_i completes; this effectively orders the transactions.
 - What if T_j already has a lock on Y and T_i later requests a lock on Y ? (Deadlock!) T_i or T_j is aborted and restarted!

Ensuring Atomicity

- DBMS ensure *atomicity* (all-or-nothing property) even if system crashes in the middle of a Xact.
- **Idea:** Keep a log (history) of all actions carried out by the DBMS while executing a set of Xacts:
 - **Before** a change is made to the database, the corresponding log entry is forced to a safe location. (Write-Ahead Log (WAL) protocol)
 - After a crash, the effects of partially executed transactions are undone using the log. (Thanks to WAL, if log entry wasn't saved before the crash, corresponding change was not applied to database!)

The Log



- The following actions are recorded in the log:
 - *Ti writes an object*: The old value and the new value.
 - Log record must go to disk before the changed page!
 - *Ti commits/aborts*: A log record indicating this action.
- Log records chained together by Xact id, so it's easy to undo a specific Xact (e.g., to resolve a deadlock).
- Log is often *archived* on “stable” storage.
- All log related activities (and in fact, all CC related activities such as lock/unlock, dealing with deadlocks etc.) are handled transparently by the DBMS.

Databases are valuable to...

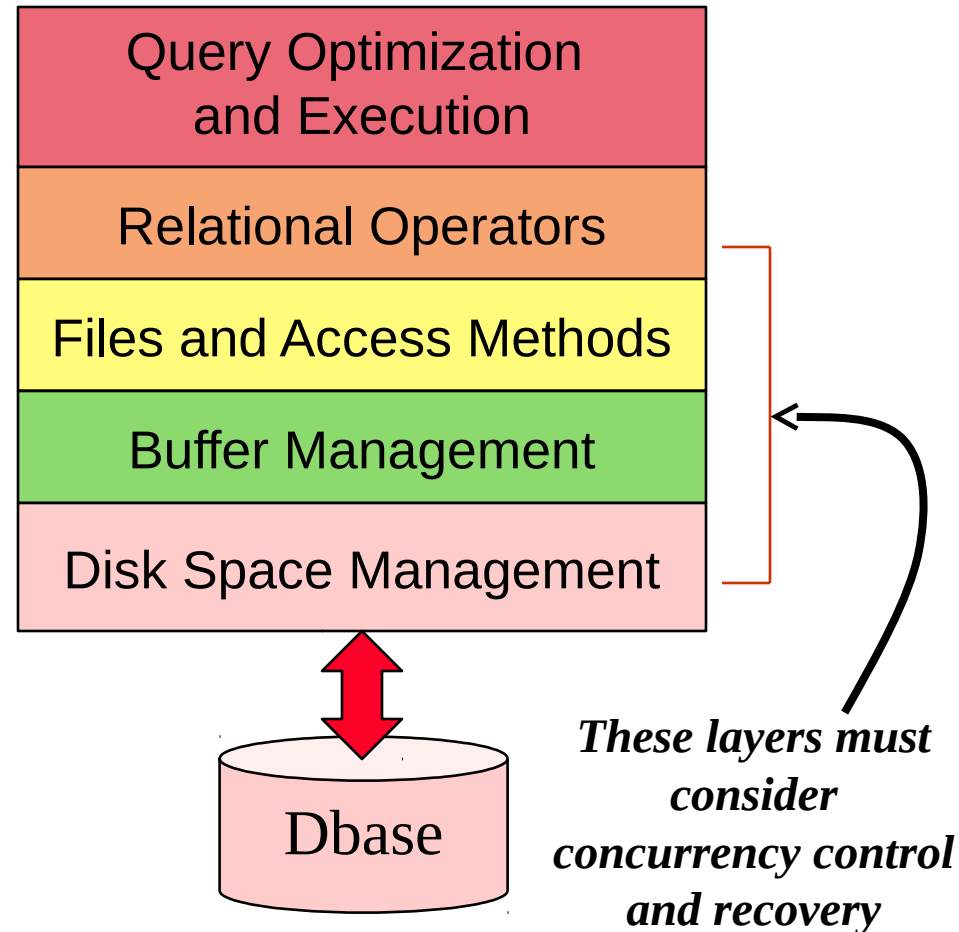
- End users (Banks, Retailers, Scientists)
- DBMS vendors (Oracle, IBM, Microsoft)
- DB application programmers
 - Makes life easier since Dbase provides guarantees
- Database administrator (DBA)
 - Designs logical/physical schemas
 - Handles security and authorization
 - Data availability, crash recovery
 - Database tuning as needs evolve



Last three must understand how a DBMS works!

Structure of a DBMS

- A typical DBMS has a layered architecture.
- The figure does not show the concurrency control and recovery components.
- This is one of several possible architectures; each system has its own variations.



Summary

- DBMS used to maintain, query large datasets.
- Benefits include recovery from system crashes, concurrent access, quick application development, data integrity, and security.
- Levels of abstraction give data independence.
- A DBMS typically has a layered architecture.
- DBMS R&D is one of the broadest, most exciting growth areas in CS.



Next

- Introduction to Relational Model
- Modeling Data
- The Entity- Relationship (ER) model

