

Chapter 7: Database Design and E-R Model

Mapping to Relational Model & Extended Features


Reduction to Relational Schemas

Reduction to Relation Schemas

- Entity sets and relationship sets can be expressed uniformly as *relation schemas* that represent the contents of the database.
- A database which conforms to an E-R diagram can be represented by a collection of schemas.
- For each entity set and relationship set there is a unique schema that is assigned the name of the corresponding entity set or relationship set.
- Each schema has a number of columns (generally corresponding to attributes), which have unique names.

Logical DB Design: ER to Relational


- Entity sets to tables:




```
CREATE TABLE Employees
(ssn CHAR(11),
name CHAR(20),
lot INTEGER,
PRIMARY KEY (ssn))
```

Representing Relationship Sets

- A many-to-many relationship set is represented as a schema with attributes for the primary keys of the two participating entity sets, and any descriptive attributes of the relationship set.
- Example: schema for relationship set *advisor*
advisor = (*s_id*, *i_id*)



Relationship Sets to Tables





Relationship Sets to Tables

- In translating a relationship set to a relation, attributes of the relation must include:
 - Keys for each participating entity set (declared as foreign keys). This set of keys is at least a *superkey* for the relation.
 - All descriptive attributes.


```
CREATE TABLE Works_In(  
    ssn CHAR(11),  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY (ssn, did),  
    FOREIGN KEY (ssn)  
        REFERENCES Employees,  
    FOREIGN KEY (did)  
        REFERENCES Departments)
```

Review: Key Constraints


- Each dept has at most one manager, according to the key constraint on Manages.




1-to-1



1-to Many



Many-to-1



Many-to-Many

Translation to relational model?

Translating ER Diagrams with Key Constraints

- Map relationship to a table:
 - Note that **did** is the key now! (did and ssn together as the key?)
 - Separate tables for Employees and Departments.
- Since each department has a unique manager, we could instead combine Manages and Departments.
 - Drawback: null value for ssn if a dept. doesn't have a manager.

Solution 1:

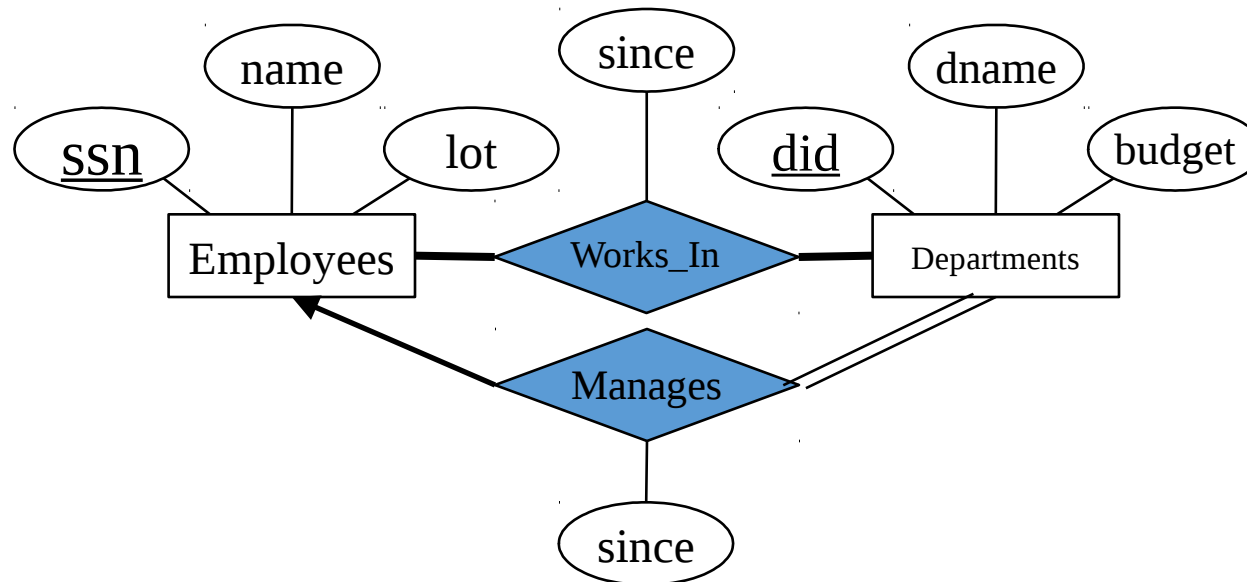
```
CREATE TABLE Manages (  
    ssn CHAR(11),  
    did INTEGER,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    FOREIGN KEY (did) REFERENCES Departments)
```

Solution 2:

```
CREATE TABLE Dept_Mgr (  
    did INTEGER,  
    dname CHAR(20),  
    budget REAL,  
    ssn CHAR(11),  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees)
```

Review: Participation Constraints

- Does every department have a manager?
 - If so, this is a *participation constraint*: the participation of Departments in Manages is said to be *total (vs. partial)*.
 - Every *did* value in Departments table must appear in a row of the Manages table (with a non-null *ssn* value)



Participation Constraints in SQL


- What is the participation constraint here?

```
CREATE TABLE Dept_Mgr (  
    did INTEGER,  
    dname CHAR(20),  
    budget REAL,  
    ssn CHAR(11) NOT NULL,  
    since DATE,  
    PRIMARY KEY (did),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE NO ACTION)
```

- What happens if we remove NOT NULL for ssn?

Review: Weak Entities

- A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.
 - Owner entity set and weak entity set must participate in a one-to-many relationship set (1 owner, many weak entities).
 - Weak entity set must have total participation in this *identifying* relationship set.



Translating Weak Entity Sets

- Weak entity set and identifying relationship set are translated into a single table.
- When the owner entity is deleted, all owned weak entities are also deleted.

```
CREATE TABLE Dep_Policy (  
    pname CHAR(20),  
    age INTEGER,  
    cost REAL,  
    ssn CHAR(11)  
    PRIMARY KEY (pname, ssn),  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE CASCADE)
```

Composite and Multivalued Attributes

<i>instructor</i>
<u>ID</u>
<i>name</i>
<i>first_name</i>
<i>middle_initial</i>
<i>last_name</i>
<i>address</i>
<i>street</i>
<i>street_number</i>
<i>street_name</i>
<i>apt_number</i>
<i>city</i>
<i>state</i>
<i>zip</i>
{ <i>phone_number</i> }
<i>date_of_birth</i>
<i>age</i> ()

- Composite attributes are flattened out by creating a separate attribute for each component attribute
 - Example: given entity set *instructor* with composite attribute *name* with component attributes *first_name* and *last_name* the schema corresponding to the entity set has two attributes *name_first_name* and *name_last_name*
 - Prefix omitted if there is no ambiguity
- Ignoring multivalued attributes, extended instructor schema is
 - *instructor*(ID, *first_name*, *middle_initial*, *last_name*, *street_number*, *street_name*, *apt_number*, *city*, *state*, *zip_code*, *date_of_birth*)


Composite and Multivalued Attributes

```
CREATE TABLE inst_phone  
(  
  ID integer,  
  Phone_number CHAR(8),  
)
```


- A multivalued attribute M of an entity E is represented by a separate schema EM
 - Schema EM has attributes corresponding to the primary key of E and an attribute corresponding to multivalued attribute M
 - Example: Multivalued attribute *phone_number* of *instructor* is represented by a schema:
inst_phone = (ID, phone_number)
 - Each value of the multivalued attribute maps to a separate tuple of the relation on schema EM
 - For example, an *instructor* entity with primary key 22222 and phone numbers 456-7890 and 123-4567 maps to two tuples:
(22222, 456-7890) and (22222, 123-4567)

Multivalued Attributes (Cont.)

- Special case: entity *time_slot* has only one attribute other than the primary-key attribute, and that attribute is multivalued
 - O m z D ' g y, j
create the one corresponding to the multivalued attribute
 - *time_slot*(*time_slot_id*, *day*, *start_time*, *end_time*)
 - Caveat: *time_slot* attribute of *section* (from *sec_time_slot*) cannot be a foreign key due to this optimization



Multivalued Attributes (Cont.)



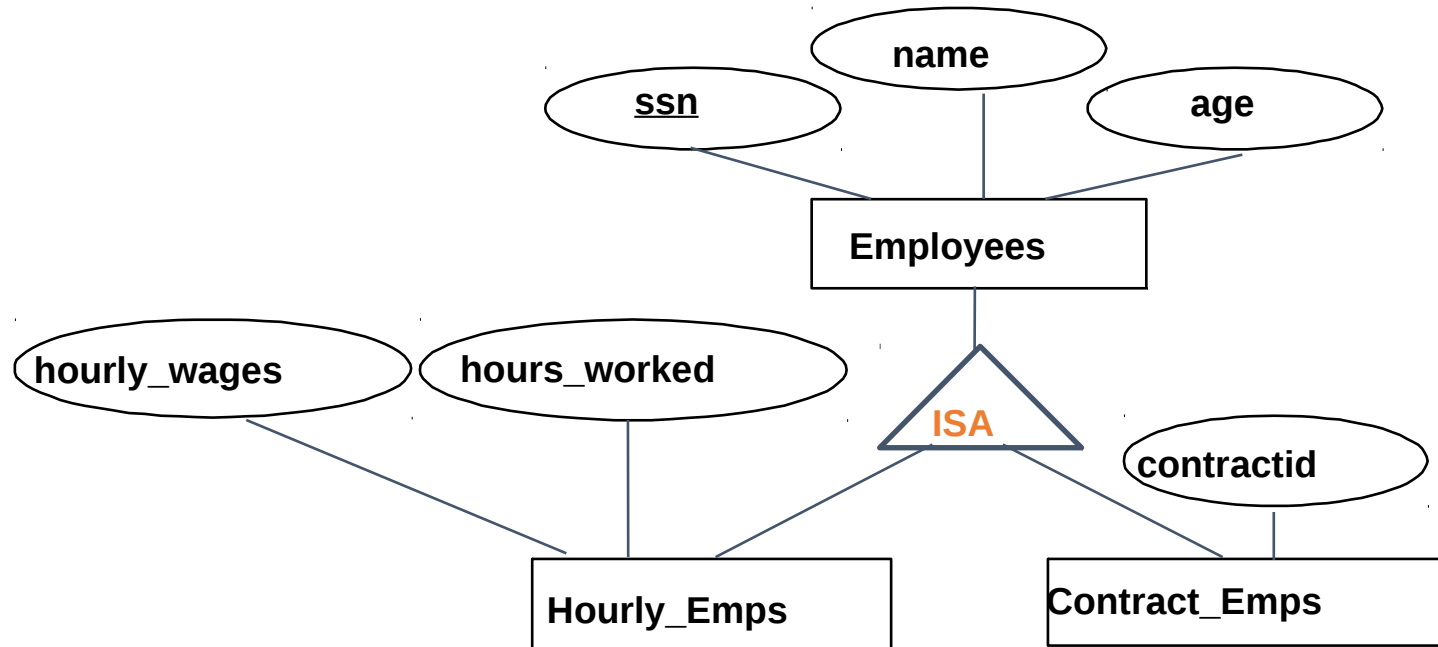
- Originally:
 - sec_time_slot: (22222, 1111)
 - Time_slot: (**1111**, {(9/5/2017,2:30,3:45), (9/7/2017,2:30,3:45)})
- Optimized:
 - sec_time_slot: (22222, 1111)
 - Time_slot: (**1111**, **9/5/2017,2:30,3:45**), (**1111**, **9/7/2017,2:30,3:45**)

Extended ER Features

Subclasses

- Sometimes, an entity set contains some entities that do share many, but not all properties with the entity set. In this case, we want to define class (entity set) hierarchies.
- A **ISA** B: every A entity is also considered to be a B entity. A *specializes* B, B *generalizes* A.
- A is called *subclass*, B is called *superclass*.
- A subclass *inherits* the attributes of a superclass, and may define additional attributes.

Subclasses



- Hourly_Emps and Contract_Emps inherit the ssn (key!), name and age attributes from Employees.
- They define additional attributes hourly_wages, hours_worked and contractid, resp.

Subclasses

- *Overlap constraints:*
Can Joe be an Hourly_Emps as well as a Contract_Emps entity?
(Hourly_Emps **OVERLAPS** Contract_Emps)
- *Disjoint constraints:*
Can Joe be an Hourly_Emps as well as a Contract_Emps entity?
(Hourly_Emps **DISJOINTS** Contract_Emps)
- *Covering constraints:*
Does every Employees entity have to be either an Hourly_Emps or a Contract_Emps entity?

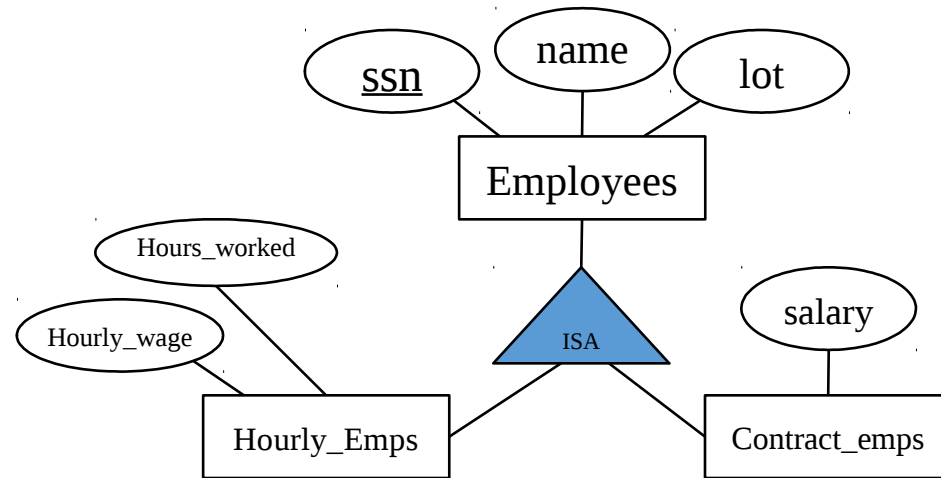
Hourly_Emps **AND** Contract_Emps **COVER** Employees
(Total Generalization vs Partial Generalization)

Subclasses

- There are several good reasons for using ISA relationships and subclasses:
 - Do not have to redefine all the attributes.
 - Can add descriptive attributes specific to a subclass.
 - To identify entity sets that participate in a relationship set as precisely as possible.
- ISA relationships form a *tree structure* (taxonomy) with one entity set serving as *root*.

Review: ISA Hierarchies

- ❖ It is often useful to subdivide entities into classes, like in an OOL
- ❖ If we declare A **ISA** B, every A entity is also considered to be a B entity.



- *Overlap constraints*: Can Joe be an Hourly_Emps as well as a Contract_Emps entity? (*Allowed/disallowed*)
- *Covering constraints*: Does every Employees entity also have to be an Hourly_Emps or a Contract_Emps entity? (*Yes/no*)

Translating ISA Hierarchies to Relations

- **General approach:**

- 3 relations: Employees, Hourly_Emps and Contract_Emps.
 - *Hourly_Emps*: Every employee is recorded in Employees. For hourly emps, extra info recorded in Hourly_Emps (*hourly_wages*, *hours_worked*, *ssn*); must delete Hourly_Emps tuple if referenced Employees tuple is deleted).
 - Queries involving all employees easy, those involving just Hourly_Emps require a join to get some attributes.

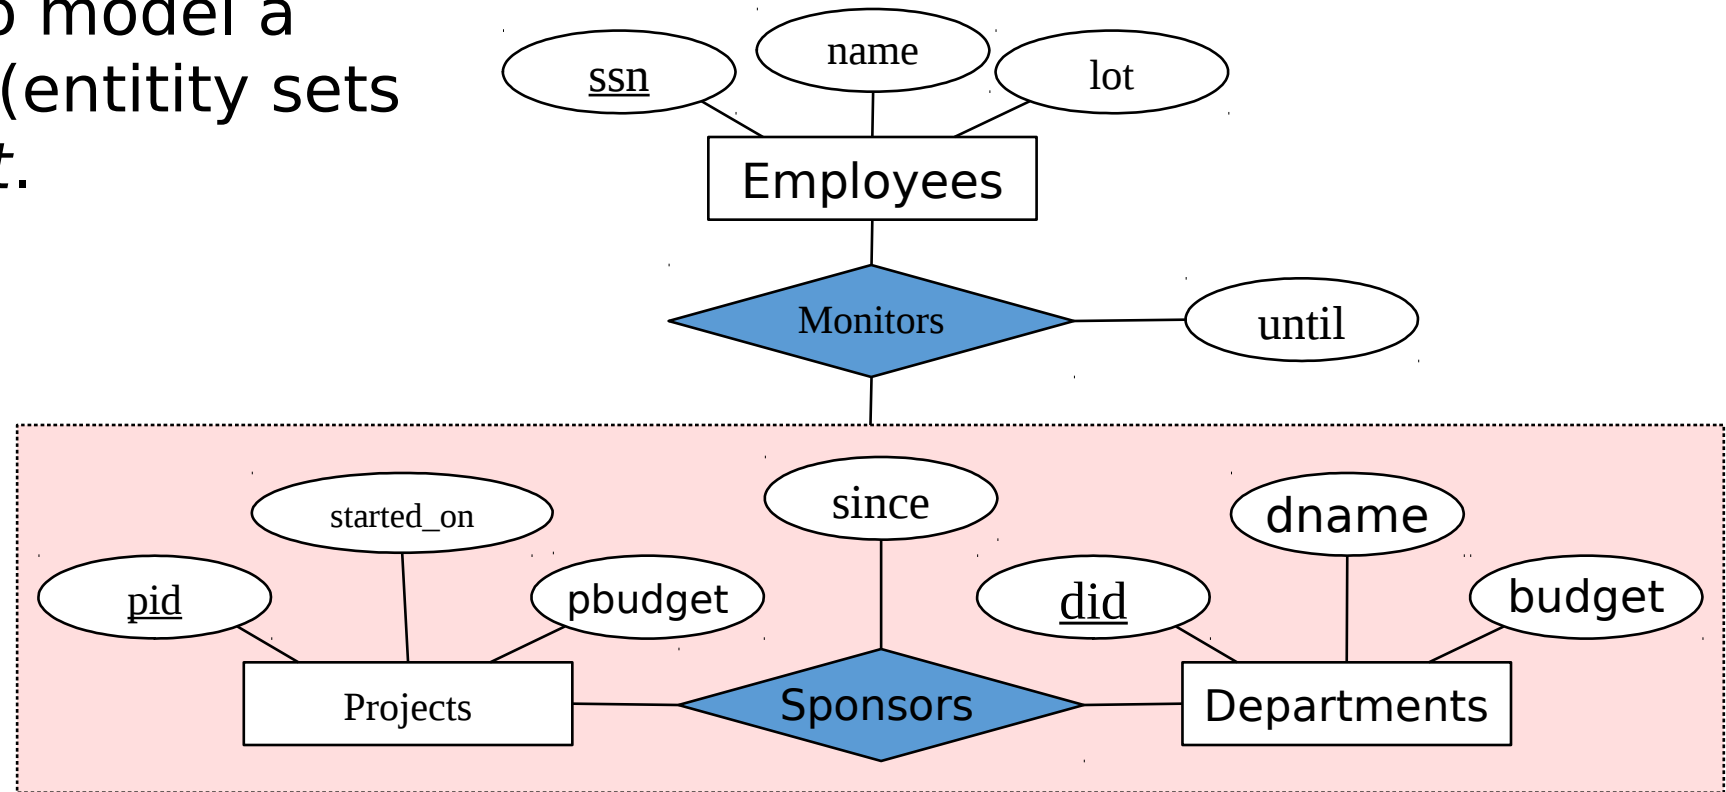
- **Alternative: Just Hourly_Emps and Contract_Emps.**

- *Hourly_Emps*: *ssn*, *name*, *lot*, *hourly_wages*, *hours_worked*.
- Each employee must be in one of these two subclasses.
- *Querying Employees requires querying both tables*

Aggregation

- Used when we have to model a relationship involving (entity sets and) a *relationship set*.

- Aggregation* allows a relationship set to be treated as an entity set for purposes of participation in (other) relationships.



- Aggregation vs. ternary relationship:*
 - Monitors is a distinct relationship, with a descriptive attribute (**until**).
 - Each sponsorship is monitored by at most one employee.

Schemas Corresponding to Aggregation

- To represent aggregation, create a schema containing
 - primary key of the aggregated relationship,
 - the primary key of the associated entity set
 - any descriptive attributes

Basically treat the aggregation relationship set as a high level entity.

Projects (pid, started_on, pbudget)

Departments (did, dname, budget)

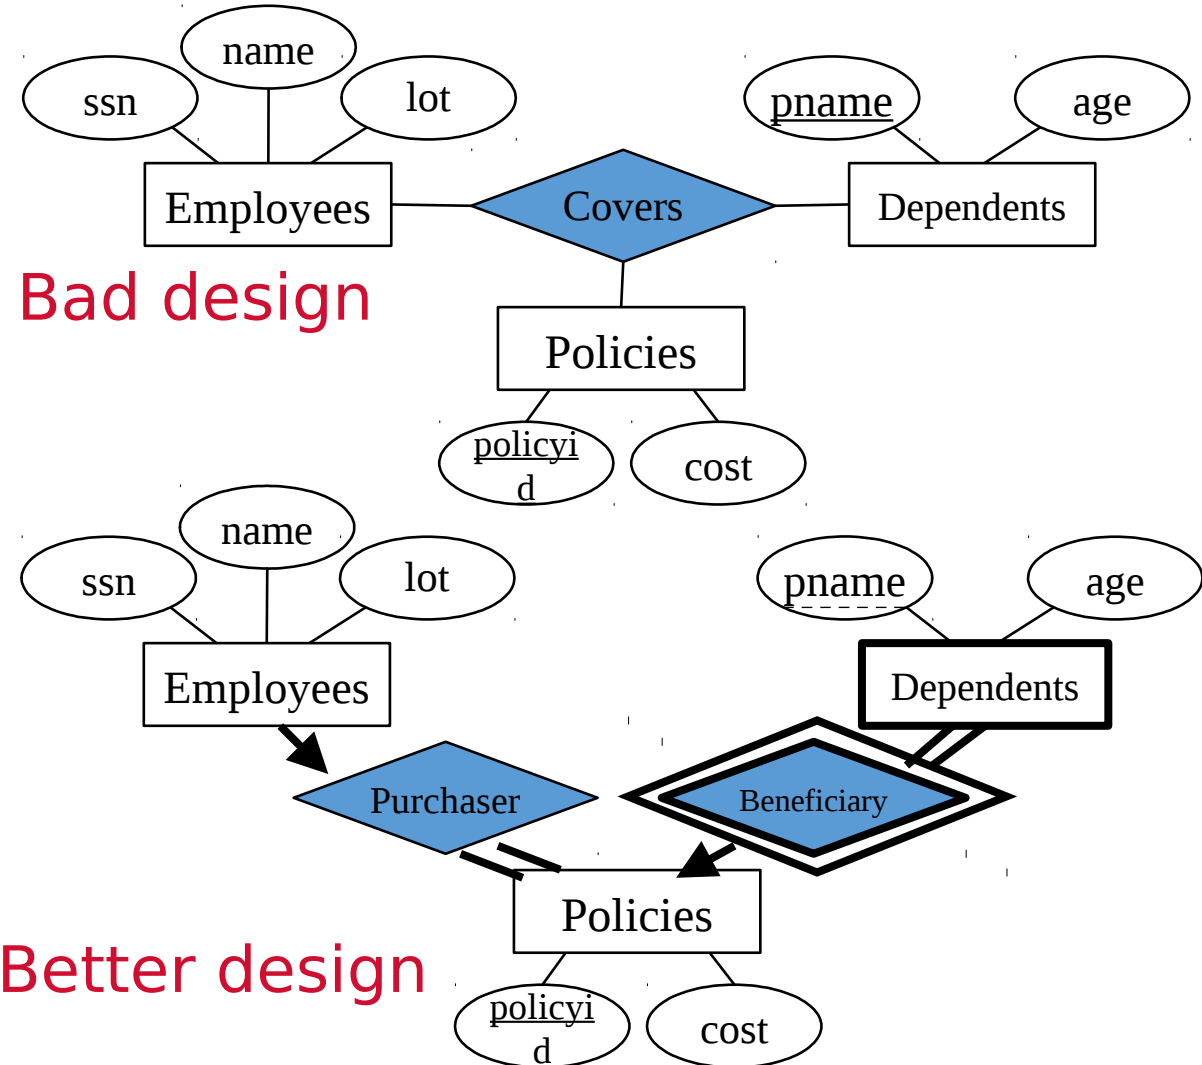
Sponsors (pid, did, since)

Employees (ssn, name, lot)

Monitors (ssn, pid, did, until)

Review: Binary vs. Ternary Relationships

- Recall what were the additional constraints implied by the the better design?
- Better:**
 - The same policy cannot be owned jointly by two or more employees.
 - Every policy must be owned by some employee.
 - Dependents is a weak entity set, and each dependent entity is uniquely identified by taking pname in conjunction with the policyid of a policy entity (which, intuitively, covers the given dependent).



Binary vs. Ternary Relationships (Contd.)

- Key constraints allow us to combine Purchaser with Policies, and Beneficiary with Dependents.
- Participation constraints lead to **NOT NULL** constraints.
- What if Policies is a weak entity set? (generic policy numbers)
 - Make ssn part of the primary key for policies

```
CREATE TABLE Policies (  
    policyid INTEGER,  
    cost REAL,  
    ssn CHAR(11) NOT NULL,  
    PRIMARY KEY (policyid).  
    FOREIGN KEY (ssn) REFERENCES Employees,  
    ON DELETE CASCADE)
```

```
CREATE TABLE Dependents (  
    pname CHAR(20),  
    age INTEGER,  
    policyid INTEGER,  
    PRIMARY KEY (pname, policyid).  
    FOREIGN KEY (policyid) REFERENCES Policies,  
    ON DELETE CASCADE)
```

Views

- A view is just a relation, but we store a *definition*, rather than a set of tuples.

```
CREATE VIEW YoungActiveStudents (login, grade)
AS SELECT S.login, E.grade
FROM Students S, Enrolled E
WHERE S.sid = E.sid and S.age<21
```

- Views can be dropped using the `DROP VIEW` command.
 - How to handle `DROP TABLE` if there's a view on the table?
 - `DROP TABLE` command has options to let the user specify this.

Views to support ISA relations

- The common elements of an ISA hierarchy can be supported using views.
- For example, consider this implementation for the employee, hourly employee, and contract employee example.

```
CREATE VIEW Employee(ssn, name, lot)
AS SELECT H.ssn, H.name, H.lot
FROM Hourly_Emps H
UNION
SELECT C.ssn, C.name, C.lot
FROM Contract_Emps C
```

Views and Security

- Views can be used to present necessary information (or a summary), while hiding details in underlying relation(s).

login	grade
smith@cs	C
smith@cs	B
smith@math	A
jones@cs	B

Design Principles

- Faithfulness
 - Design must be faithful to the specification / reality.
 - Relevant aspects of reality must be represented in the model.
- Avoiding redundancy
 - Redundant representation blows up ER diagram and makes it harder to understand.
 - Redundant representation wastes storage.
 - Redundancy may lead to inconsistencies in the database.

Design Principles

- Keep it simple
 - The simpler, the easier to understand for some (external) reader of the ER diagrams.
 - Avoid introducing more elements than necessary.
 - If possible, prefer attributes over entity sets and relationship sets.
- Formulate constraints as far as possible
 - A lot of data semantics can (and should) be captured.
 - But some constraints cannot be captured in ER diagrams.

High-Level Design With E-R Model


- Major E-R design decisions
 - Should a concept be modeled as an entity or an attribute?
 - Should a concept be modeled as an entity or a relationship?
 - What relationships to use: binary or ternary?
 - The use of a strong or weak entity set.
 - The use of specialization/generalization – contributes to modularity in the design.
 - The use of aggregation – can treat the aggregate entity set as a single unit without concern for the details of its internal structure.

Entity Sets vs. Attributes


- Should *address* be an attribute of Employees or an entity (connected to Employees by a relationship)?
- Depends upon the use we want to make of address information, and the semantics of the data:
 - If we have several addresses per employee, *address* must be an entity (since attributes cannot be set-valued).
 - If the structure (city, street, etc.) is important, e.g., we want to retrieve employees in a given city, *address* must be modeled as an entity (since attribute values are atomic).

Entity Sets vs. Attributes

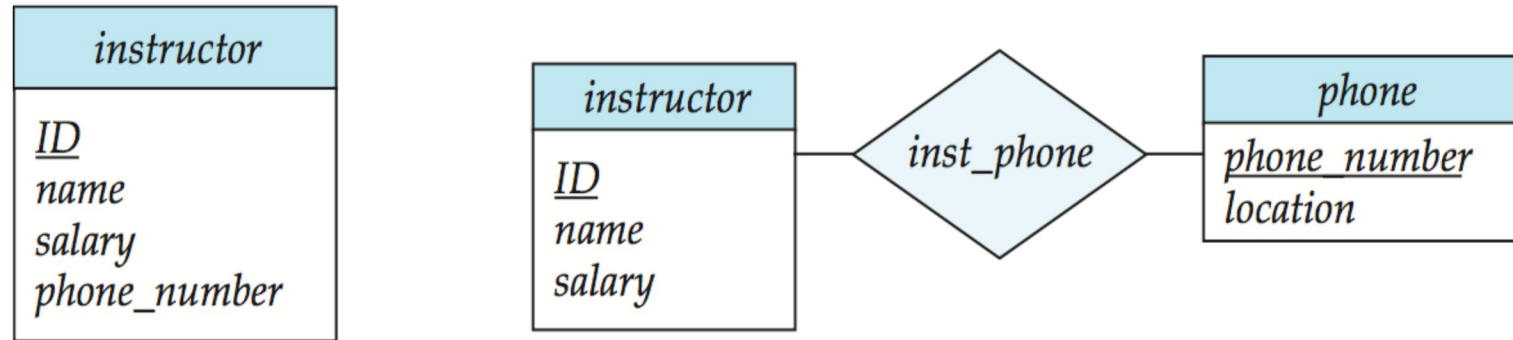
- Works_In2 does not allow an employee to work in the same department for two or more periods (why?).



- We want to record *several values of the descriptive attributes for each instance of this relationship.*



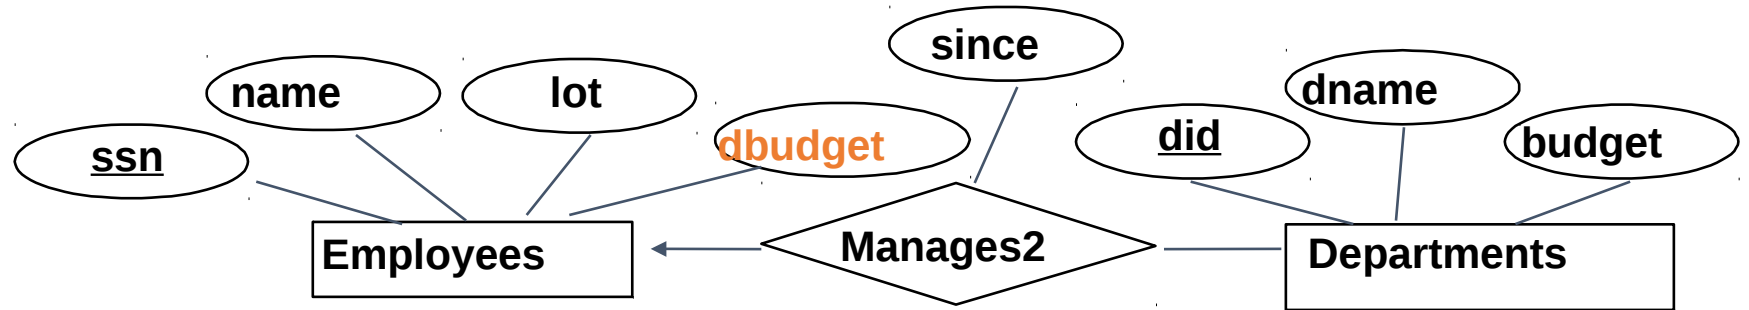
Entity Sets vs. Attributes



- Use of phone as an entity allows extra information about phone numbers (plus multiple phone numbers)

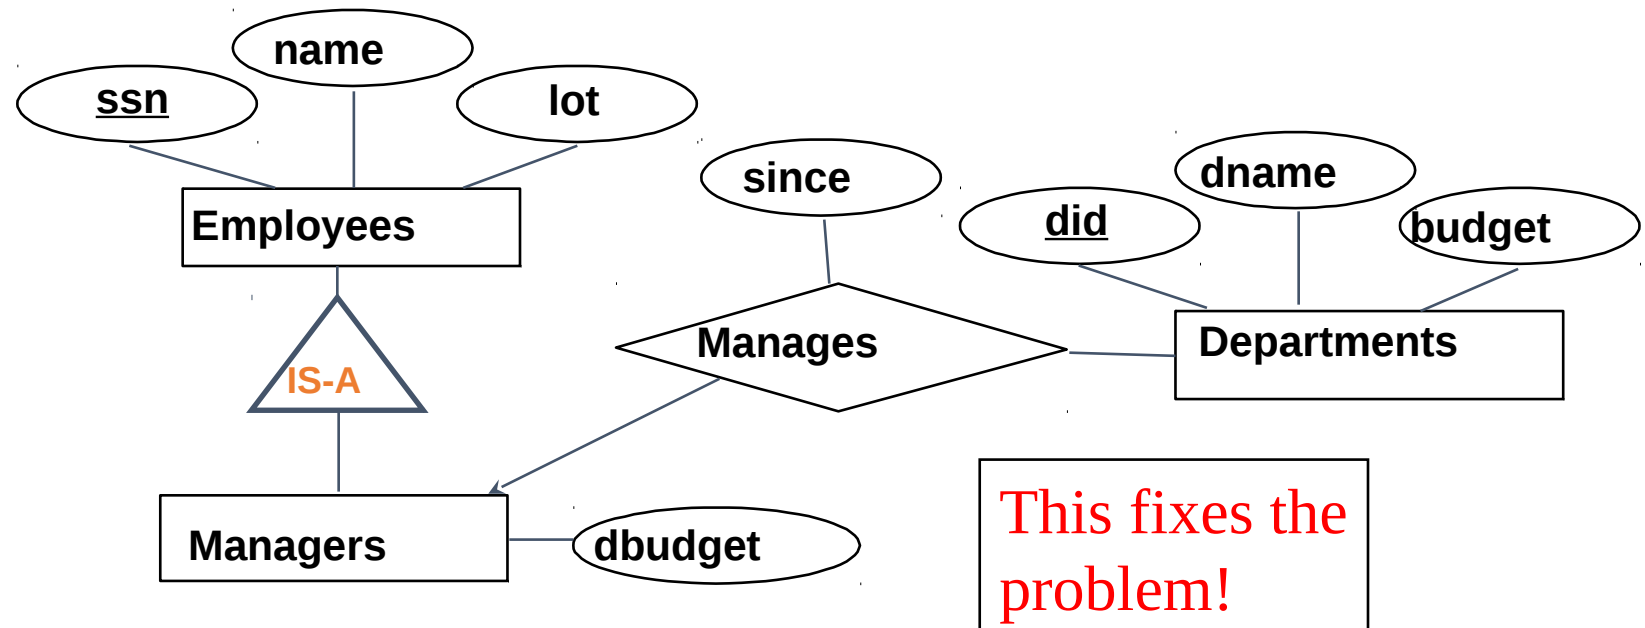
Entity vs. Relationship

- What about this diagram?




Not every employee is a manager!!

- The following ER diagram is more appropriate and avoids the above problems!



This fixes the problem!

Entity vs. Relationship



- This ER diagram o.k. if a manager gets a *separate* discretionary budget for each dept.
- But what if a manager gets a discretionary budget that covers *all* managed depts?
 - **Redundancy** of *dbudget*, which is stored for each dept managed by the manager.
 - **Misleading**: suggests *dbudget* tied to managed dept.

Summary of Conceptual Design

- *Conceptual design follows requirements analysis,*
 - Yields a high-level description of data to be stored
- ER model popular for conceptual design
 - Constructs are expressive, close to the way people think about their applications.
- Basic constructs: *entities, relationships, and attributes* (of entities and relationships).
- Some additional constructs: *weak entities, ISA hierarchies, and aggregation.*
- Note: There are many variations on ER model.

Summary of ER (Contd.)

- Several kinds of integrity constraints can be expressed in the ER model: *key constraints*, *participation constraints*, and *overlap/covering constraints* for ISA hierarchies. Some *foreign key constraints* are also implicit in the definition of a relationship set.
 - Some constraints (notably, *functional dependencies*) cannot be expressed in the ER model.
 - Constraints play an important role in determining the best database design for an enterprise.

Summary of ER (Contd.)

- ER design is *subjective*. There are often many ways to model a given scenario! Analyzing alternatives can be tricky, especially for a large enterprise. Common choices include:
 - Entity vs. attribute, entity vs. relationship, binary or n-ary relationship, whether or not to use ISA hierarchies, and whether or not to use aggregation.
- Ensuring good database design: resulting relational schema should be analyzed and refined further. FD information and normalization techniques are especially useful.

Next:

- Introduction to SQL