

# Chapter 3: Introduction to SQL (continued)

# Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*
  - Example:  $5 + \text{null}$  returns null
- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.

```
select name  
from instructor  
where salary is null
```

# Null Values and Three Valued Logic

- Any comparison with *null* returns *unknown*
  - Example:  $5 < null$  or  $null <> null$  or  $null = null$
- Three-valued logic using the truth value *unknown*:
  - OR:  $(unknown \textbf{ or } true) = true$ ,  
 $(unknown \textbf{ or } false) = unknown$   
 $(unknown \textbf{ or } unknown) = unknown$
  - AND:  $(true \textbf{ and } unknown) = unknown$ ,  
 $(false \textbf{ and } unknown) = false$ ,  
 $(unknown \textbf{ and } unknown) = unknown$
  - NOT:  $(\textbf{not } unknown) = unknown$
  - “*P is unknown*” evaluates to true if predicate *P* evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*

# Three-Valued Logic

- To understand how AND, OR, and NOT work in 3-valued logic, think of TRUE = 1, FALSE = 0, and UNKNOWN =  $\frac{1}{2}$ .
- AND = MIN; OR = MAX, NOT( $x$ ) =  $1-x$ .

Example:

$$\begin{aligned}\text{TRUE AND (FALSE OR NOT(UNKNOWN))} &= \\ \text{MIN(1, MAX(0, (1 - } \frac{1}{2} \text{ )))} &= \\ \text{MIN(1, MAX(0, } \frac{1}{2} \text{ ))} &= \text{MIN(1, } \frac{1}{2} \text{ )} = \frac{1}{2}.\end{aligned}$$

# Surprising Example


- From the following Sells relation:

| bar       | beer | price |
|-----------|------|-------|
| Joe's Bar | Bud  | NULL  |

SELECT bar

FROM Sells

WHERE price < 2.00 OR price >= 2.00;



## Reason: 2-Valued Laws $\neq$ 3-Valued Laws

- Some common laws, like commutativity of AND, hold in 3-valued logic.
  - **Unknown** AND False AND True = False AND True AND **Unknown**
- But not others,
- e.g., the *law of the excluded middle* :  $p \text{ OR NOT } p = \text{TRUE}$ .
  - When  $p = \text{UNKNOWN}$ ,  
 $\text{MAX}( \frac{1}{2}, (1 - \frac{1}{2}) ) = \frac{1}{2} \neq 1$ .

# Aggregate Functions

- These functions operate on the multiset of values of a column of a relation, and return a value

**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values

# Aggregate Functions (Cont.)

- Find the average salary of instructors in the Computer Science department
  - **select avg** (*salary*)  
**from** *instructor*  
**where** *dept\_name*= 'Comp. Sci.';
- Find the total number of instructors who teach a course in the Spring 2015 semester
  - **select count** (**distinct** *ID*)  
**from** *teaches*  
**where** *semester* = 'Spring' **and** *year* = 2015
- Find the number of tuples in the *course* relation
  - **select count** (\*)  
**from** *course*;



# Aggregate Functions - Group By

- Find the average salary of instructors in each department

```
select dept_name, avg (salary)
from instructor
group by dept_name;
```

Note: departments with no instructor will not appear in result

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 76766     | Crick       | Biology          | 72000         |
| 45565     | Katz        | Comp. Sci.       | 75000         |
| 10101     | Srinivasan  | Comp. Sci.       | 65000         |
| 83821     | Brandt      | Comp. Sci.       | 92000         |
| 98345     | Kim         | Elec. Eng.       | 80000         |
| 12121     | Wu          | Finance          | 90000         |
| 76543     | Singh       | Finance          | 80000         |
| 32343     | El Said     | History          | 60000         |
| 58583     | Califieri   | History          | 62000         |
| 15151     | Mozart      | Music            | 40000         |
| 33456     | Gold        | Physics          | 87000         |
| 22222     | Einstein    | Physics          | 95000         |

| <i>dept_name</i> | <i>avg_salary</i> |
|------------------|-------------------|
| Biology          | 72000             |
| Comp. Sci.       | 77333             |
| Elec. Eng.       | 80000             |
| Finance          | 85000             |
| History          | 61000             |
| Music            | 40000             |
| Physics          | 91000             |

# Queries With GROUP BY and HAVING

```
SELECT      [DISTINCT]
target-list
FROM        relation-list
WHERE       qualification
GROUP BY    grouping-list
HAVING      group-qualification
```

- The *target-list* contains (i) attribute names (ii) terms with aggregate operations (e.g., MIN (*S.age*)).
- The attribute names in (i) above must be a subset of *grouping-list*. Intuitively, each answer tuple corresponds to a *group*, and these attributes **must have a single value per group**.  
(A *group* is a set of tuples that have the same value for all attributes in *grouping-list*.)

# Interpretation Semantics (Conceptual Evaluation Strategy)

- ***Interpretation Semantics*** (different than the ***Execution Semantics***):
  - Compute the cross-product of ***relation-list***.
  - Discard resulting tuples if they fail ***qualifications***.
  - Delete attributes that are not in ***target-list***.
  - If **DISTINCT** is specified, eliminate duplicate rows.
- This strategy is the least efficient way to compute a query! An optimizer will find more efficient strategies to compute ***the same answer***.

# Grouping

- We may follow a SELECT-FROM-WHERE expression by GROUP BY and a list of attributes.
- The relation that results from the SELECT-FROM-WHERE is grouped according to the values of all those attributes, and any aggregation is applied only within each group.

# Aggregation (Cont.)

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list

Example: “Find average salary of instructors for each department. Instructor (ID, name, dept\_name, salary)  
/\* erroneous query \*/

```
select dept_name, ID, avg (salary)
from instructor
group by dept_name;
```

What is the correct SQL query?

Remove ID

**Query:** “Find the number of instructors in each department who teach a course in the Fall 2009 semester.”

Instructor (ID, name, dept\_name, salary)

Course (course\_id, title, dept\_name, credits)

Section (course\_id, section\_id, semester, year)

Teaches (ID, course\_id, sec\_id, semester, year)

Department(dept\_name, building, budget)

In SQL?

**Query:** “Find the number of instructors in each department who teach a course in the Fall 2009 semester.”

Instructor (ID, name, dept\_name, salary)

Course (course\_id, title, dept\_name, credits)

Section (course\_id, section\_id, semester, year)

Teaches (ID, course\_id, sec\_id, semester, year)

Department(dept\_name, building, budget)

```
SELECT I.dept_name, Count (distinct T.ID)
```

```
FROM Instructor I, Teaches T
```

```
WHERE I.ID= T.ID
```

```
    AND T.semester=“Fall”
```

```
    AND T.year= “2009”
```

```
GROUP BY I.dept_name
```

**Query: “For each department, find the number of instructors who teach a course offered by that department in the Fall 2009 semester.”**

Instructor (ID, name, dept\_name, salary)

Course (course\_id, title, dept\_name, credits)

Section (course\_id, section\_id, semester, year)

Teaches (ID, course\_id, sec\_id, semester, year)

Department(dept\_name, building, budget)

```
SELECT D.dept_name, Count (T.ID)
```

```
FROM Department D, Teaches T, COURSE C
```

```
WHERE D.dept_name = C.dept_name // A course offered by the department
```

```
    AND C.course_id = T.course_id
```

```
    AND T.semester='Fall'
```

```
    AND T.year= '2009'
```

```
GROUP BY D.dept_name
```

Is this SQL query correct?

**Missing distinct!**



**Query: “For each department, find the number of instructors who teach a course offered by that department in the Fall 2009 semester.”**

Instructor (ID, name, dept\_name, salary)

Course (course\_id, title, dept\_name, credits)

Section (course\_id, section\_id, semester, year)

Teaches (ID, course\_id, sec\_id, semester, year)

Department(dept\_name, building, budget)

```
SELECT I.dept_name, Count (distinct T.ID)
FROM Instructor I, Teaches T, COURSE C
WHERE I.dept_name = C.dept_name // A course offered by the department
      AND C.course_id = T.course_id
      AND T.semester='Fall'
      AND T.year= '2009'
GROUP BY I.dept_name
```

# Aggregate Functions - Having Clause

- Find the names and average salaries of all departments whose average salary is greater than 60000

```
60000  select dept_name, avg (salary)
      from instructor
      group by dept_name
      having avg (salary) > 60000;
```

**Note:** predicates in the **having** clause are applied **after** the formation of groups whereas predicates in the **where** clause are applied **before** forming groups

| <i>dept_name</i> | <i>avg(salary)</i> |
|------------------|--------------------|
| Physics          | 91000              |
| Elec. Eng.       | 80000              |
| Finance          | 85000              |
| Comp. Sci.       | 77333              |
| Biology          | 72000              |
| History          | 61000              |

The result relation for the query “Find the average salary of instructors in those departments where the average salary is more than \$60,000.”

# Null Values and Aggregates

- Total all salaries

```
select sum (salary )  
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount
- All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
  - count returns 0
  - all other aggregates return null

## Next:

More SQL, Nested Subqueries, Correlated subqueries, ...