

# Chapter 5: Advanced SQL

- Accessing SQL From a Programming Language
  - Dynamic SQL
    - JDBC and ODBC
  - Embedded SQL
- SQL Data Types and Schemas
- Functions and Procedural Constructs
- Triggers

# Accessing SQL from a general Purpose programming language

Two main reasons:

1. Not all queries can be expressed in SQL
2. Nondeclarative Actions (e.g., printing, output formatting, interactive actions, ...)

SQL commands can be called from within a host language (e.g., C++ or Java) program.

- SQL statements can refer to *host variables* (including special variables used to return status).
- Must include a statement to *connect* to the right database.

## Two Main approaches:

- **Dynamic SQL:** allows program to construct an SQL query as a char. String at run time, submit the query and retrieve results tuple-at-time into program variables.
- **Embedded SQL:** SQL statements are identified at compile time by preprocessor, complied and optimized by database system, and replaced by appropriate code and function calls before programming language compiler is invoked.

## Impedance mismatch:

SQL relations are (multi-) sets of records, with no *a priori* bound on the number of records. Sets are not supported cleanly in procedural programming languages

- SQL supports a mechanism called a [cursor](#) to handle this.

# Embedded SQL

- Approach: Embed SQL in the host language.
  - An DBMS-specific preprocessor converts the SQL statements into special API calls.
  - Then a regular compiler is used to compile the code.
- Language constructs:
  - Connecting to a database:  
EXEC SQL CONNECT
  - Declaring variables:  
EXEC SQL BEGIN (END) DECLARE SECTION
  - Statements:  
EXEC SQL Statement;

# Embedded SQL: Variables

```
EXEC SQL BEGIN DECLARE SECTION  
char c_sname[20];  
long c_sid;  
short c_rating;  
float c_age;  
EXEC SQL END DECLARE SECTION
```

Two special “error” variables:

- SQLCODE (long, is negative if an error has occurred)
- SQLSTATE (char[6], predefined codes for common errors)

**Problem 1:** Data types in SQL must be recognized by host language, and vice versa.

**Sol’n:** Use casting (of data values of one type into values of the other one).

**Problem 2:** DB query language: set-oriented, and host language: value oriented.

**Sol’n:** Use cursors.

# Cursors

- Can declare a cursor on a relation or query statement (which generates a relation).
- Can *open* a cursor, and repeatedly *fetch* a tuple then *move* the cursor, until all tuples have been retrieved.
  - Can use a special clause, called **ORDER BY**, in queries that are accessed through a cursor, to control the order in which tuples are returned.
  - Fields in ORDER BY clause must also appear in SELECT clause.
  - The **ORDER BY** clause, which orders answer tuples, is *only* allowed in the context of a cursor.
- Can also modify/delete tuple pointed to by a

Cursor that gets names of students who've enrolled in EECS 341, in alphabetical order

```
EXEC SQL DECLARE sinfo CURSOR FOR  
  SELECT S.sname  
  FROM Student S, Enrolled E  
  WHERE S.sid=E.sid AND E.cid="EECS 341"  
  ORDER BY S.sname
```

```
OPEN sinfo;
```

# Embedded SQL (Cont.)

- The **open** statement causes the query to be evaluated

**EXEC SQL open c END\_EXEC**

- The **fetch** statement causes the values of one tuple in the query result to be placed on host language variables.

**EXEC SQL fetch c into :si, :sn END\_EXEC**

Repeated calls to **fetch** get successive tuples in the query result

- A variable called SQLSTATE in the SQL communication area (SQLCA) gets set to '02000' to indicate no more data is available
- The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

**EXEC SQL close c END\_EXEC**

Note: above details vary with language. For example, the Java embedding defines Java iterators to step through result tuples.



# Embedding SQL in C: An Example

```
char SQLSTATE[6];
EXEC SQL BEGIN DECLARE SECTION
char c_sname[20]; char c_major; short c_year;
EXEC SQL END DECLARE SECTION
c_major = 'CS';
EXEC SQL DECLARE sinfo CURSOR FOR
    SELECT S.sname, S.year FROM Students S
    WHERE S.major = :c_major // prefix ":" for variables defined
    ORDER BY S.sname;      // in host program
EXEC SQL OPEN sinfo
do {
    EXEC SQL FETCH sinfo INTO :c_sname, :c_year;
    printf( c_sname, c_year);
} while (SQLSTATE != '02000');
EXEC SQL CLOSE sinfo;
```

# Example Query

- From within a host language, find the ID and name of students who have completed more than the number of credits stored in variable **credit\_amount**.

- Specify the query in SQL and declare a *cursor* for it

**EXEC SQL**

**declare c cursor for**

**select** *ID, name*

**from** *student*

**where tot\_cred** > *:credit\_amount*

**END\_EXEC**

# Updates Through Cursors

- Can update tuples fetched by cursor by declaring that the cursor is for update

```
declare c cursor for  
select *  
from instructor  
where dept_name = 'Music'  
for update
```

- To update tuple at the current location of cursor *c*

```
update instructor  
set salary = salary + 100  
where current of c
```

# Dynamic SQL

- SQL query strings are not always known at compile time (e.g., spreadsheet, graphical DBMS frontend): Allows construction of SQL statements on-the-fly.

- **Example:**

```
char c_sqlstring[]=
    {"DELETE FROM Students WHERE year = "freshmen"};
EXEC SQL PREPARE readytogo FROM :c_sqlstring;
EXEC SQL EXECUTE readytogo;
```

# Database APIs: Alternative to embedding

Rather than modify compiler, add library with database calls (API).

- Special standardized interface: procedures/objects
- Passes SQL strings from language; presents result sets in a language-friendly way.
- ODBC (open database connectivity)
- Sun's *JDBC*: Java API
- Supposedly DBMS-neutral
  - a “driver” traps the calls and translates them into DBMS-specific code.
  - database can be across a network.

# JDBC: Architecture

- Four architectural components:
  - Application (initiates and terminates connections, submits SQL statements).
  - Driver manager (loads JDBC driver).
  - Driver (connects to data source, transmits requests and returns/translates results and error codes).
  - Data source (processes SQL statements).

# JDBC Classes and Interfaces

Steps to submit a database query:

1. Load the JDBC driver
2. Connect to the data source
3. Execute SQL statements

# Executing SQL Statements in JDBC

- Three different ways of executing SQL statements:
  - Statement (both static and dynamic SQL statements)
  - PreparedStatement (semi-static SQL statements)
  - CallableStatement (stored procedures)
- PreparedStatement class:
  - Precompiled, parameterized SQL statements
  - Structure is fixed.
  - Values of parameters are determined at run-time.



# ResultSet in JDBC

- PreparedStatement.executeUpdate only returns the number of affected records.
- PreparedStatement.executeQuery returns data, encapsulated in a ResultSet object (a cursor).

```
ResultSet rs=pstmt.executeQuery(sql);
```

```
// rs is now a cursor
```

```
While (rs.next()) {
```

```
    // process the data
```

```
}
```

# ResultSet (Contd.)

A ResultSet is a very powerful cursor:

- `previous()`: moves one row back.
- `absolute(int num)`: moves to the row with the specified number.
- `relative (int num)`: moves forward or backward.
- `first()` and `last()`

# JDBC: Exceptions and Warnings

- Most of java.sql can throw a SQLException if an error occurs.
- SQLWarning is a subclass of SQLException; not as severe (they are not thrown and their existence has to be explicitly tested)

# SQL Procedures

- The *dept\_count* function could instead be written as procedure:

```
create procedure dept_count_proc (in dept_name  
varchar(20),  
  
out d_count  
integer)  
begin  
    select count(*) into d_count  
    from instructor  
    where instructor.dept_name =  
    dept_count_proc.dept_name  
end
```

- Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.

```
declare d_count integer;  
call dept_count_proc( 'Physics', d_count);
```